


## CpSc 4620/6620: Database Management Systems (DBMS) (TEXNH Approach)


# Relational Algebra

James Wang




## Relational Algebra

- ✿ Relational Algebra builds the foundation for SQL Queries.
- ✿ Operations:
  - ✿ Selection:  $\sigma_p(R)$ .  $p$  - predicates,  $R$  - relation
  - ✿ Projection:  $\pi_a(R)$ .  $a$  - attributes,  $R$  - relation
  - ✿ Union:  $R \cup S$ .  $R, S$  - relations.
  - ✿ Intersection:  $R \cap S$ .  $R, S$  - relations.
  - ✿ Difference:  $R - S$ .  $R, S$  - relations.
  - ✿ Cross-Product:  $R \times S$ .  $R, S$  - relations.
  - ✿ Renaming:  $\rho_x(E)$ .  $E$  is a relational expression.
  - ✿ Join:  $R \bowtie_p S$ .  $p$  - predicates,  $R, S$  - relations
  - ✿ Division:  $R \div S$ .  $R, S$  - relations
  - ✿ More ...




## Primitive Operations

- ✿ The six primitive operators of relational algebra are the *selection*, the *projection*, the *Cartesian product* (also called the *cross product* or *cross join*), the *set union*, the *set difference*, and the *rename*.
- ✿ Constraints:
  - ✿ For set union and set difference, the two relations involved must be union-compatible, that is, the two relations must have the same set of attributes.
  - ✿ As set intersection can be defined in terms of set difference, the two relations involved in set intersection must also be union-compatible.
  - ✿  $R \times S$  is defined as follows:  
 $R \times S = \{r \cup s \mid r \in R, s \in S\}$   
Note: two relations involved in cross product must not have a common attribute name, otherwise, the result won't be a relation.



## Set Union and Difference

- ✿  $R \cup S$  is defined as follows:  
 $R \cup S = \{x \mid x \in R \text{ or } x \in S\}$ , where  $R$  and  $S$  are relations with the same attributes.
- ✿  $R - S$  (same times  $R \setminus S$ ) is defined as follows:  
 $R - S = \{x \mid x \in R \text{ and } x \notin S\}$ , where  $R$  and  $S$  are relations with the same attributes.
- ✿ Set intersection  $R \cap S$  can be defined by set difference.  
 $R \cap S = \{x \mid x \in R \text{ and } x \in S\}$   
 $= R - (R - S)$



## PROJECTION

- ✿ In relational algebra, a projection is a unary operation written as  $\pi_{a_1, \dots, a_n}(R)$  where  $a_1, \dots, a_n$  is a set of attribute names.
- ✿ The result of such projection is defined as the set that is obtained when all tuples in  $R$  are restricted to the set  $\{a_1, \dots, a_n\}$ .
- ✿ Example:
 


Name	Age	Weight
Harry	34	80
Sally	28	64
George	29	70
Helena	54	54
Peter	34	80

Age	Weight
34	80
28	64
29	70
54	54

$\pi_{\text{Age, Weight}}(\text{Person})$

DISTINCT

(http://en.wikipedia.org/wiki/Projection\_%28relational\_algebra%29)



## SELECTION

- ✿ In relational algebra, a selection is a unary operation written as  $\sigma_{a\theta b}(R)$  or  $\sigma_{a\theta v}(R)$  where:
  - ✿  $a$  and  $b$  are attribute names
  - ✿  $\theta$  is a binary operation in the set
  - ✿  $v$  is a value constant
  - ✿  $R$  is a relation
- ✿ Selection is sometimes called a restriction to avoid confusion with SQL's use of SELECT.
- ✿ The selection  $\sigma_{a\theta b}(R)$  selects all those tuples in  $R$  for which  $\theta$  holds between the  $a$  and the  $b$  attribute.
- ✿ The selection  $\sigma_{a\theta v}(R)$  selects all those tuples in  $R$  for which  $\theta$  holds between the  $a$  attribute and the value  $v$ .

## SELECTION Examples

**Person**

Name	Age	Weight
Harry	34	80
Sally	28	64
George	29	70
Helena	54	54
Peter	34	80

$\sigma_{\text{Age} \geq 34}(\text{Person})$

Name	Age	Weight
Harry	34	80
Helena	54	54
Peter	34	80

$\sigma_{\text{Age} = \text{Weight}}(\text{Person})$

Name	Age	Weight
Helena	54	54

([http://en.wikipedia.org/wiki/Selection\\_\(relational\\_algebra\)](http://en.wikipedia.org/wiki/Selection_(relational_algebra)))

## RENAME

- In relational algebra, a rename is a unary operation written as  $\rho_{a/b}(R)$  where:
  - $a$  and  $b$  are attribute names
  - $R$  is a relation
- The result is identical to  $R$  except that the  $b$  field in all tuples is renamed to an  $a$  field.
- Example:**

Name	EmployeeID
Harry	3415
Sally	3345

EmployeeName	EmployeeID
Harry	3415
Sally	3345

## Natural join

- Natural join is a binary operator that is written as  $R \bowtie S$  where  $R$  and  $S$  are relations. The result of the natural join is the set of all combinations of tuples in  $R$  and  $S$  that are equal on their common attribute names.
- Example:**

Name	EmpId	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales

DeptName	Manager
Finance	George
Sales	Harriet
Production	Charles

$\text{Employee} \bowtie \text{Dept}$

Name	EmpId	DeptName	Manager
Harry	3415	Finance	George
Sally	2241	Sales	Harriet
George	3401	Finance	George
Harriet	2202	Sales	Harriet

## $\theta$ -join and equijoin

- If we want to combine tuples from two relations where the combination condition is not simply the equality of shared attributes,  $\theta$ -join (or theta-join) is necessary.
- The  $\theta$ -join is a binary operator that is written as  $R \bowtie_{a \theta b} S$  or  $R \bowtie_{a \theta v} S$ , where  $a$  and  $b$  are attribute names,  $\theta$  is a binary relation in the set  $\{<, \leq, =, >, \geq\}$ ,  $v$  is a value constant, and  $R$  and  $S$  are relations.
- The result of this operation consists of all combinations of tuples in  $R$  and  $S$  that satisfy the relation  $\theta$ . The result of the  $\theta$ -join is defined only if the headers of  $S$  and  $R$  are disjoint, that is, do not contain a common attribute.
- In case the operator  $\theta$  is the equality operator ( $=$ ) then this join is also called an equijoin.

## Example of $\theta$ -join

- Consider tables *Car* and *Boat* which list models of cars and boats and their respective prices. Suppose a customer wants to buy a car and a boat, but she doesn't want to spend more money for the boat than for the car. The  $\theta$ -join on the relation  $\text{CarPrice} \geq \text{BoatPrice}$  produces a table with all the possible options.

CarModel	CarPrice
CarA	20'000
CarB	30'000
CarC	50'000

BoatModel	BoatPrice
Boat1	10'000
Boat2	40'000
Boat3	60'000

$\text{Car} \bowtie_{\text{CarPrice} \geq \text{BoatPrice}} \text{Boat}$

CarModel	CarPrice	BoatModel	BoatPrice
CarA	20'000	Boat1	10'000
CarB	30'000	Boat1	10'000
CarC	50'000	Boat1	10'000
CarC	50'000	Boat2	40'000

## Semijoin


- The semijoin is joining similar to the natural join and written as  $R \ltimes S$  where  $R$  and  $S$  are relations. The result of the semijoin is only the set of all tuples in  $R$  for which there is a tuple in  $S$  that is equal on their common attribute names.
- Example:**

Name	EmpId	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Production

DeptName	Manager
Sales	Harriet
Production	Charles

$\text{Employee} \ltimes \text{Dept}$

Name	EmpId	DeptName
Sally	2241	Sales
Harriet	2202	Production



## Antijoin


The antijoin, written as  $R \bowtie S$  where  $R$  and  $S$  are relations, is similar to the natural join, but the result of an antijoin is only those tuples in  $R$  for which there is NOT a tuple in  $S$  that is equal on their common attribute names.

**Example:**

Employee		
Name	EmpId	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Production

Dept	
DeptName	Manager
Sales	Harriet
Production	Charles

Employee $\bowtie$ Dept		
Name	EmpId	DeptName
Harry	3415	Finance
George	3401	Finance



## Division


The division is a binary operation that is written as  $R \div S$ . The result consists of the restrictions of tuples in  $R$  to the attribute names unique to  $R$ , i.e., in the header of  $R$  but not in the header of  $S$ , for which it holds that all their combinations with tuples in  $S$  are present in  $R$ .

**Example:**

Completed	
Student	Task
Fred	Database1
Fred	Database2
Fred	Compiler1
Eugene	Database1
Eugene	Compiler1
Sara	Database1
Sara	Database2

DBProject	
Task	
Database1	
Database2	

Completed $\div$ DBProject	
Student	
Fred	
Sara	



## Left outer join


The left outer join is written as  $R \ltimes S$  where  $R$  and  $S$  are relations. The result of the left outer join is the set of all combinations of tuples in  $R$  and  $S$  that are equal on their common attribute names, in addition (loosely speaking) to tuples in  $R$  that have no matching tuples in  $S$ .

**Example:**

Employee		
Name	EmpId	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales
Tim	1123	Executive

Dept	
DeptName	Manager
Sales	Harriet
Production	Charles

Employee $\ltimes$ Dept			
Name	EmpId	DeptName	Manager
Harry	3415	Finance	NULL
Sally	2241	Sales	Harriet
George	3401	Finance	NULL
Harriet	2202	Sales	Harriet
Tim	1123	Executive	NULL



## Right outer join


The right outer join is written as  $R \rtimes S$  where  $R$  and  $S$  are relations. The result of the right outer join is the set of all combinations of tuples in  $R$  and  $S$  that are equal on their common attribute names, in addition to tuples in  $S$  that have no matching tuples in  $R$ .

**Example:**

Employee		
Name	EmpId	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales
Tim	1123	Executive

Dept	
DeptName	Manager
Sales	Harriet
Production	Charles

Employee $\rtimes$ Dept			
Name	EmpId	DeptName	Manager
Sally	2241	Sales	Harriet
Harriet	2202	Sales	Harriet
NULL	NULL	Production	Charles



## Outer join (Full outer join)


The full outer join is written as  $R \ltimes S$  where  $R$  and  $S$  are relations. The result of the full outer join is the set of all combinations of tuples in  $R$  and  $S$  that are equal on their common attribute names, in addition to tuples in  $R$  that have no matching tuples in  $S$  and tuples in  $S$  that have no matching tuples in  $R$  in their common attribute names.

**Example:**

Employee		
Name	EmpId	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales
Tim	1123	Executive

Dept	
DeptName	Manager
Sales	Harriet
Production	Charles

Employee $\ltimes$ Dept			
Name	EmpId	DeptName	Manager
Harry	3415	Finance	NULL
Sally	2241	Sales	Harriet
George	3401	Finance	NULL
Harriet	2202	Sales	Harriet
Tim	1123	Executive	NULL
NULL	NULL	Production	Charles



## References

- An Introduction to Database Systems, Eighth Edition, C. J. Date, Addison Wesley, 2004, ISBN: 0-321-19784-4.
- [http://en.wikipedia.org/wiki/Relational\\_algebra](http://en.wikipedia.org/wiki/Relational_algebra)